

## Design Decisions

### Streaming Instrumentation?

WASM allows streaming parsing and compilation, i.e., start while only parts of the module have been transmitted over the network (useful for browsers, especially on mobile).

Pros:

- not possible for x86 but for WASM
- can execute parts even before instrumentation is complete, i.e., lower instrumentation latency
- new “term”

Cons:

- implementation: either in-engine or as a proxy
- ??? only useful if instrumentation is faster than network speed, otherwise the latter catches up anyway
- actual use in practice? ??? a developer will have the modules local anyway, not many will be doing instrumentation in production, especially not on mobile

### Parallel Instrumentation? (related to Streaming)

Sections in WASM modules can be decoded/compiled independently of each other, so parallelize instrumentation per section/function.

Pros:

- not possible for x86 but for WASM
- higher instrumentation bandwidth
- new “term”
- unlike streaming instrumentation, can be implemented externally / as a standalone “WASM module transformer”

Cons:

- ??? multithreading is hard to implement (?)

MP: How about single-pass, modular instrumentation? I.e., it could be used for streaming instrumentation and it could be easily parallelized.

### Record/Replay?

As per its paper, Jalangi does this, i.e., it first records a trace of all memory accesses/used values and later uses them in the replay phase for shadow execution (i.e., the actual analysis on top of the instrumentation).

Pros:

- useful for time-traveling debugging?
- potentially the recording is quicker than very slow analyses, so the instrumented binary would run quicker than with “online” analysis

Cons:

- ??? Jalangi2 doesn't use this anymore (why?)
- at least doubles the memory requirements of the instrumented binary cf. the uninstrumented one

- is the analysis really more expensive than recording every memory access? (I suspect not, since the analysis results are probably much less data to store/process.) So in reality, does this slow down the instrumented binary more than “online” analysis?

MP: Maybe as a second step, once we have an instrumentation tool?

## Static vs. Dynamic (JIT) Instrumentation

Dynamic: instrument only “on-demand” when some code is executed (== jumped to), not everything beforehand. PIN and Valgrind do it.

Pros Dynamic:

- necessary for x86, since static disassembly is hard / impossible (linear sweep is not reliable because a) code and data are mixed and b) variable-length instructions / very dense ISA)
- only instruments the parts that actually get executed
- change instrumentation at runtime (e.g., add more specific instrumentation at runtime or disable after some time)

Pros Static / Cons Dynamic:

- robust disassembly in WASM is possible (it is designed exactly for that)
- JIT machinery is complicated:
  - needs instrumented code cache HashMap + some quick lookup before that
  - need to include instrumentation framework at runtime of the instrumented binary
  - needs to patch code at runtime, i.e., change jump targets from the “dispatcher” to the newly instrumented code (this is not even possible in WASM!?)
- ??? most of a binary is probably executed, thus dynamic instrumentation only delays the instrumentation but has to do it eventually anyway (TODO: evaluate! is this true?)

MP: Sounds like static makes more sense for WASM.

## Implementation of the Instrumentation: In-Browser/In-Engine vs. External Program

Pros In-Browser:

- can reuse existing parts of the codebase, e.g., WASM decoder (?)
- necessary for streaming instrumentation

Pros External:

- easier to get going: no need to modify the engine, recompile browser everytime, read and find around their source code
- we cannot compete with Browser vendors anyway, so rather not try
- users do not need to install custom browser to use instrumentation
- we do not need to track changes in the browser
- works across different browser and different vendors

MP: If implemented as an external tool, couldn’t we compile it to WASM and then run in the browser anyway? I’d go for external.

## Disassemble & Resynthesize vs. Copy & Annotate

I think Valgrind paper has coined these terms: D&R means having your own IR to which the binary is parsed first, then the instrumentation added in this IR, and then serialized to a binary again. In particular, the original binary is discarded completely and everything only works on the IR. C&A implies that the original binary's code is copied over and instrumentation inserted in between. The distinction doesn't seem 100% clear to me, especially since PIN (which the Valgrind paper says is C&A) also rewrites original binary instructions (they do register reallocation).

The distinction in WASM is even less clear: WASM is executed on a VM/JIT anyway.

Pros D&R:

- original code and instrumentation code are optimized together, e.g., calls in both cases could be inlined by the JIT during execution
- for x86: less register spilling, i.e., the analysis code can use registers just as good as the original code, whereas in C&A the analysis code must make sure not to modify any registers of the original binary -> lots of spills
- for WASM: there is already a well-specified, small IR (unlike for x86, where this is a major challenge to correctly capture all of the ISA's behavior)
- for WASM: parsing binary to an IR and serializing again should be easy to make round trip (unlike for x86)

Pros C&A:

- for x86: more robust, less danger of changing semantics of an instruction if the original instructions are executed unchanged
- for x86: less implementation effort, no need to capture the huge CISC instruction set.

MP: We should use an IR, either an existing one or our own (the latter might also become a contribution).

## “Heavy” vs “Light” Instrumentation

Valgrind paper argues their framework enables analyses that are not possible with “lightweight” PIN/DynamoRIO, e.g., Memcheck (TODO: what does it actually do?). Where is the boundary between lightweight and heavyweight?

Pros Heavy:

- more complex analyses possible

Cons Heavy:

- performance
- larger API to implement (?)

MP: We should aim for a general framework that supports many different, possibly complex analyses. Performance definitely also matters, but should be a secondary goal.

## API Design: based on Event Streams (and Filters) vs. Callbacks vs. ???

RoadRunner paper talks a lot about their API design: Analyses are implemented as “filters”, i.e., methods that capture some events (generated by the framework, e.g., `FieldAccessEvent`) and give the uncaptured events through to later “filters” in the “tool chain”.

MP: We could also have a low-level API, which provides access to basic events, and a high-level API, which summarizes low-level events into higher-level events. To make more concrete suggestions, I'll have to learn more about WASM.

## Language Choice: Instrumentation Framework

C++:

- already exists some tooling (WABT)
- fast (?)
- I am not an expert in C++ :(

Rust:

- lots of small community for WASM already: a parser/unparser, a small interpreter, maybe Mozillas internal implementation (?)
- easier/safer for multithreading stuff than C++
- I already know the language a bit (and I love it) :)

JavaScript:

- huge community, close to the web world
- very high level
- possibly slow
- no chance of every integrating that inside the engines of the browsers :(

## Language Choice: User Analyses

WASM:

- easy to insert in the existing binary then :)
- too low-level to implement large analyses in it

JavaScript:

- how expensive is the JS-WASM interop? If it is expensive, instrumentation could become too slow
- JavaScript itself is already slow

C++, Rust, any language that compiles to WASM:

- fast, but still high-level

MP: Can't we compile JS to WASM? Any language that compiles to WASM seems fine.

## Other

- WASM has only structured control flow: is this something new/noteworthy?

## Possible Analyses / Use-Cases for Instrumentation

*Simple* ones, i.e., analysis itself is well known

- instruction counting: which instruction is executed how often
  1. per whole module (simplest case)
  2. per function
- basic block counting / coverage
  - needs some control-flow analysis
- taint analysis
  - simplest version: just 1 bit, “tainted” or not
  - needs “shadow values”, i.e., associated meta information for each memory “cell” (byte?)
  - needs “shadow execution”, i.e., computation on this meta information, invoked when computation on actual values happens
- memory access tracing
  - record address of each read/write
  - show regions that are never accessed (maybe as an image?) -> should have not been allocated?
  - show regions that are accessed very often -> potential for optimization?
- call counting
  1. which function is called how often
    - need to handle direct and indirect calls
  2. from which other function (build flame graph?)
- null/empty instrumentation: run instrumentation framework but do no user analysis

Known but not as simple to implement:

- general time profiling: which function is executed how long, flamegraph?
  - precise (actually measure time) vs. sampling (periodically observe which is currently executing, is this possible in WASM?)
- fault injection: flip bits/change bytes in WASM memory
- concolic execution (?)
- track origin of null values: when null value is read, where was it written from?

New analyses, possibly *specific to WebAssembly*:

- allocation profiler: invocations of `Memory.grow(num_pages)`
  - which function (as a flamegraph)
  - how often
  - with what page count
- JavaScript-WASM interop linter
  - correctness: do JavaScript types and WebAssembly types (mis)match? (do other instrumentation first, maybe I learn about interop in the process anyway)
  - performance: how long does the conversion take?
- data type mismatches in WebAssembly: detect pointer arguments and remember how often how many bytes are read from this ptr, if different sizes -> possible data type mismatch?

Other:

- record & replay, for time-traveling debugging